

# Structured Query Language (SQL)

- SQL is the most widely used commercial relational database language

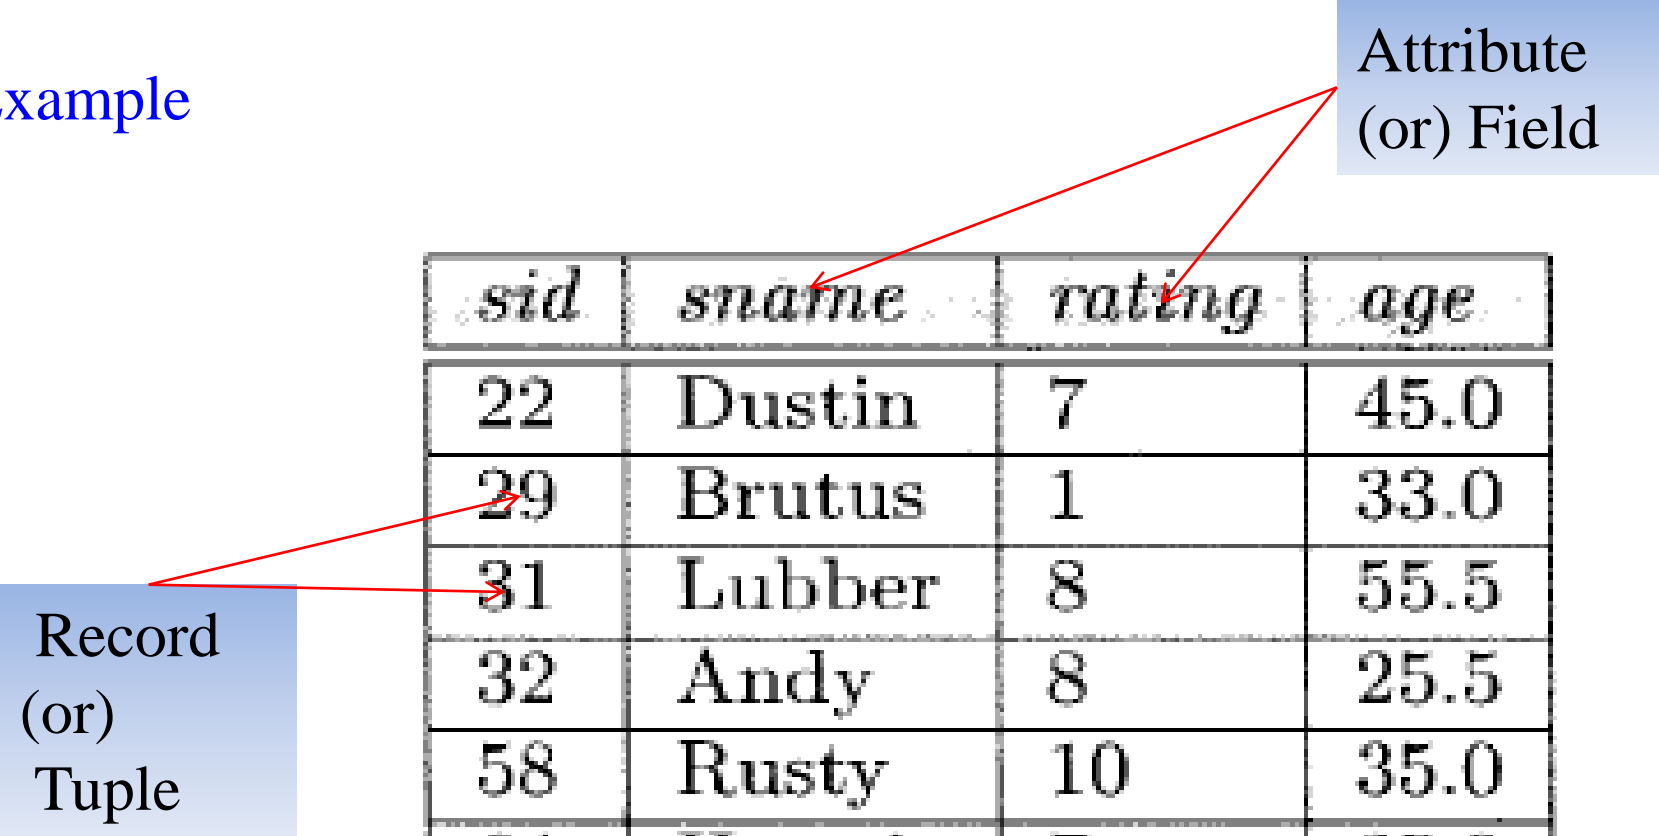
## **THE FORM OF A BASIC SQL QUERY**

SELECT [DISTINCT] field names  
FROM table names  
WHERE condition

- SELECT clause contains fields to be displayed in the result
- FROM clause contains table names
- Optional WHERE clause contains conditions on the tables mentioned in the FROM clause

## Example

Attribute  
(or) Field



| <i>sid</i> | <i>sname</i> | <i>rating</i> | <i>age</i> |
|------------|--------------|---------------|------------|
| 22         | Dustin       | 7             | 45.0       |
| 29         | Brutus       | 1             | 33.0       |
| 31         | Lubber       | 8             | 55.5       |
| 32         | Andy         | 8             | 25.5       |
| 58         | Rusty        | 10            | 35.0       |
| 64         | Horatio      | 7             | 35.0       |
| 71         | Zorba        | 10            | 16.0       |
| 74         | Horatio      | 9             | 35.0       |
| 85         | Art          | 3             | 25.5       |
| 95         | Bob          | 3             | 63.5       |

Record  
(or)  
Tuple

Fig 4.1 Sailors table

*Question: Find the names and ages of all sailors*

Query: SELECT **DISTINCT** sname, age  
FROM Sailors

Output:

| <i>sname</i> | <i>age</i> |
|--------------|------------|
| Dustin       | 45.0       |
| Brutus       | 33.0       |
| Lubber       | 55.5       |
| Andy         | 25.5       |
| Rusty        | 35.0       |
| Horatio      | 35.0       |
| Zorba        | 16.0       |
| Art          | 25.5       |
| Bob          | 63.5       |

**Note:** The DISTINCT keyword can be used to return only distinct (different) values from the specific field

*Question: Find the names and ages of all sailors*

Query:       SELECT sname, age  
              FROM Sailors

Output:

| <i>sname</i> | <i>age</i> |
|--------------|------------|
| Dustin       | 45.0       |
| Brutus       | 33.0       |
| Lubber       | 55.5       |
| Andy         | 25.5       |
| Rusty        | 35.0       |
| Horatio      | 35.0       |
| Zorba        | 16.0       |
| Horatio      | 35.0       |
| Art          | 25.5       |
| Bob          | 63.5       |

*Question: Find all sailors with a rating above 7*

Query: SELECT sid, sname, rating, age  
FROM Sailors  
WHERE rating > 7

(or)  
SELECT \*  
FROM Sailors  
WHERE rating > 7

Output:

|    |         |    |      |
|----|---------|----|------|
| 31 | Lubber  | 8  | 55.5 |
| 32 | Andy    | 8  | 25.5 |
| 58 | Rusty   | 10 | 35.0 |
| 71 | Zorba   | 10 | 16.0 |
| 74 | Horatio | 9  | 35.0 |

## AND, OR and NOT Operators

*Question: Find the names of sailors who have reserved boat number 103*

Query: SELECT sname  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid AND R.bid=103

Output:

| SNAME   |
|---------|
| Dustin  |
| Lubber  |
| Horatio |

## BETWEEN

- Used to define range limits

*Question: Find all sailors whose age is in between 45.0 and 63.5*

Query: **SELECT \***  
**FROM Sailors**  
**WHERE age BETWEEN 45.0 AND 63.5**

Output:

| <b>SID</b> | <b>SNAME</b> | <b>RATING</b> | <b>AGE</b> |
|------------|--------------|---------------|------------|
| 22         | Dustin       | 7             | 45         |
| 31         | Lubber       | 8             | 55.5       |
| 95         | Bob          | 3             | 63.5       |



# IN

- Used to check whether an attribute value matches a value contained within a set of listed values

*Question: Find all sailors whose age is in the list of values(15.0,33.2,45.7,63.5)*

Query: SELECT \*  
FROM Sailors  
WHERE age IN (15.0,33.2,45.7,63.5)

Output:

| SID | SNAME | RATING | AGE  |
|-----|-------|--------|------|
| 95  | Bob   | 3      | 63.5 |

## STRING operators

- “%” character is used to match any substring
- “\_” character is used to match any character
- Expresses patterns by using the ‘like’ comparison operator

### Example1

```
SELECT *  
FROM Sailors  
WHERE sname LIKE '_u%'
```

Output:

| <b>SID</b> | <b>SNAME</b> | <b>RATING</b> | <b>AGE</b> |
|------------|--------------|---------------|------------|
| 22         | Dustin       | 7             | 45         |
| 31         | Lubber       | 8             | 55.5       |
| 58         | Rusty        | 10            | 35         |

Example2

```
SELECT *  
FROM Sailors  
WHERE sname LIKE 'A_d_'
```

Output:

| <b>SID</b> | <b>SNAME</b> | <b>RATING</b> | <b>AGE</b> |
|------------|--------------|---------------|------------|
| 32         | Andy         | 8             | 25.5       |

## SET operators

- Operations such as *union*, *intersect*, *minus* and *exists* operate on relations
- Corresponding to relational-algebra operations  $U$ ,  $\cap$  and  $-$
- Relations participating in the operations must be **compatible**; i.e., must have same set of attributes

<query 1> <set operator> <query 2>

- ***union*** returns a table consisting of all rows either appearing in the result of <query 1> or in the result of <query 2>

Example (union)

```
SELECT *  
FROM Sailors  
  
UNION  
  
SELECT *  
FROM Sailors1
```

Output:

| <i>sid</i> | <i>sname</i> | <i>rating</i> | <i>age</i> |
|------------|--------------|---------------|------------|
| 22         | Dustin       | 7             | 45.0       |
| 29         | Brutus       | 1             | 33.0       |
| 31         | Lubber       | 8             | 55.5       |
| 32         | Andy         | 8             | 25.5       |
| 58         | Rusty        | 10            | 35.0       |
| 64         | Horatio      | 7             | 35.0       |
| 71         | Zorba        | 10            | 16.0       |
| 74         | Horatio      | 9             | 35.0       |
| 85         | Art          | 3             | 25.5       |
| 95         | Bob          | 3             | 63.5       |

Example (intersect)

```
SELECT *  
FROM Sailors  
INTERSECT  
SELECT *  
FROM Sailors1
```

Output:

| <i>sid</i> | <i>sname</i> | <i>rating</i> | <i>age</i> |
|------------|--------------|---------------|------------|
| 22         | Dustin       | 7             | 45.0       |
| 29         | Brutus       | 1             | 33.0       |
| 31         | Lubber       | 8             | 55.5       |
| 32         | Andy         | 8             | 25.5       |
| 58         | Rusty        | 10            | 35.0       |
| 64         | Horatio      | 7             | 35.0       |
| 71         | Zorba        | 10            | 16.0       |
| 74         | Horatio      | 9             | 35.0       |
| 85         | Art          | 3             | 25.5       |
| 95         | Bob          | 3             | 63.5       |

Example (minus)

```
SELECT *  
FROM Sailors  
MINUS  
SELECT *  
FROM Sailors1
```

Output: no rows selected

## Nested Queries

- A nested query is a query that has another query embedded within it
- The embedded query is called a subquery

- A subquery typically appears within the WHERE clause of a query
- Subqueries can sometimes appear in the FROM clause or the HAVING clause
- In the nested queries, the inner subquery is completely independent of the outer query

## Introduction to Nested Queries

*Question: Find the names of sailors who have reserved boat 103*

*Query:*

```
SELECT S.sname
```



```
FROM Sailors S
WHERE S.sid IN
(SELECT R.sid FROM Reserves R WHERE R.bid=103)
```

Output:

| SNAME   |
|---------|
| Dustin  |
| Lubber  |
| Horatio |

*Question: Find the names of sailors who have reserved a blue boat*

*Query:*

```
SELECT S.sname
FROM Sailors S
```

```
WHERE S.sid IN  
(SELECT R.sid FROM Reserves R WHERE R.bid IN  
(SELECT B.bid FROM Boats B WHERE B.color='blue'))
```

Output:

| SNAME   |
|---------|
| Dustin  |
| Horatio |

## Correlated Nested Queries

- In Correlated Nested Queries, inner subquery could depend on the row that is currently being examined in the outer query

*Question: Find the names of sailors who have reserved boat 103*

*Query:*

```
SELECT S.sname  
FROM Sailors S  
WHERE EXISTS  
(SELECT * FROM Reserves R WHERE R.bid = 103 AND  
R.sid = S.sid)
```

*Output:*

| <b>SNAME</b> |
|--------------|
| Dustin       |
| Lubber       |
| Horatio      |

- The EXISTS operator is another set comparison operator, such as IN
- It allows us to **test whether a set is nonempty**. Thus, for each Sailor row  $S$ , we test whether the set of Reserves rows  $R$  such that  $R.bid = 103$  AND  $S.sid = R.sid$  is nonempty. If so, sailor  $S$  has reserved boat 103, and we retrieve the name
- The subquery clearly depends on the current row  $S$  and must be re-evaluated for each row in Sailors
- The occurrence of  $S$  in the subquery (in the form of the literal  $S.sid$ ) is called a **correlation**, and such queries are called **correlated queries**

## COMPARISION OPERATORS

- These operators can be used in 'WHERE' clause and 'HAVING' clause

| <b>SYMBOL</b>  | <b>MEANING</b>           |
|----------------|--------------------------|
| =              | Equal to                 |
| <              | Less than                |
| <=             | Less than or equal to    |
| >              | Grater than              |
| >=             | Greater than or equal to |
| <> or != or ^= | Not equal to             |

### *Example*

*Question: Find sailors whose rating is better than some sailor called Horatio*

*Query:*

```
SELECT S1.sname, S1.rating  
FROM Sailors S1  
WHERE S1.rating > ANY (SELECT S2.rating FROM  
Sailors S2 WHERE S2.sname='Horatio' )
```

*Output:*

| SNAME   | RATING |
|---------|--------|
| Rusty   | 10     |
| Zorba   | 10     |
| Horatio | 9      |
| Lubber  | 8      |
| Andy    | 8      |

*Question: Find sailors whose rating is better than some sailor called Horatio*

*Query:*

```
SELECT S1.sname, S1.rating
FROM Sailors S1
WHERE S1.rating > ALL ( SELECT S2.rating FROM
Sailors S2 WHERE S2.sname='Horatio' )
```

*Output:*

| SNAME RATING |    |
|--------------|----|
| Rusty        | 10 |
| Zorba        | 10 |

## AGGREGATE OPERATORS

- In addition to simply retrieving data, we often want to perform some computation or summarization
- SQL supports the following aggregate operators which can be applied on any column, say A, of a relation(table):
  1. COUNT ([DISTINCT] A): The number of (unique) values in the A column
  2. SUM ([DISTINCT] A): The sum of all (unique) values in the A column
  3. AVG ([DISTINCT] A): The average of all (unique) values in the A column



4. MAX (A): The maximum value in the A column

5. MIN (A): The minimum value in the A column

**Note:** not specify DISTINCT in conjunction with MIN or MAX

**Examples:**

*Question: Find the average age of all sailors*

*Query:*

```
SELECT AVG (age)
```

```
FROM Sailors
```

*Output:*

| AVG(AGE) |
|----------|
| 36.9     |

*Question: Find the name and age of the oldest sailor*

*Query:*

```
SELECT S1.sname, S1.age
FROM Sailors S1
WHERE S1.age = ( SELECT MAX (S2.age) FROM Sailors
S2 )
```

*Output:*

| SNAME | AGE  |
|-------|------|
| Bob   | 63.5 |

*Question: Count the number of sailors*

*Query:*

```
SELECT COUNT (*)
FROM Sailors
```

*Output:*

| COUNT(*) |
|----------|
| 10       |

## The GROUP BY and HAVING Clauses

- We have applied aggregate operators to all (qualifying) rows in a relation(table)
- GROUP BY used to apply aggregate operators to each of a number of groups of rows in a relation
- HAVING is used to place a condition, which is applied on the groups of rows

general form:

SELECT [DISTINCT] fieldname

FROM table names

WHERE condition

GROUP BY fieldname

HAVING group-condition

## Examples

*Question: Find the number of sailors belongs to each rating level*

*Query:*

```
SELECT rating, COUNT(rating)
FROM Sailors
GROUP BY rating
```

*Output:*

| <b>RATING</b> | <b>COUNT(RATING)</b> |
|---------------|----------------------|
| 1             | 1                    |
| 3             | 2                    |
| 7             | 2                    |
| 8             | 2                    |
| 9             | 1                    |
| 10            | 2                    |

*Question: Find the age of the youngest sailor for each rating level*

*Query:*

```
SELECT rating, MIN (age)
FROM Sailors
GROUP BY rating
```

*Output:*

| <b>RATING</b> | <b>MIN(AGE)</b> |
|---------------|-----------------|
| 1             | 33              |
| 3             | 25.5            |
| 7             | 35              |
| 8             | 25.5            |
| 9             | 35              |
| 10            | 16              |

*Question: Find the age of the youngest sailor for each rating level, which is greater than 7*

*Query:*

```
SELECT rating, MIN(age)
FROM Sailors
GROUP BY rating
HAVING rating>7
```

*Output:*

| <b>RATING</b> | <b>MIN(AGE)</b> |
|---------------|-----------------|
| 8             | 25.5            |
| 9             | 35              |
| 10            | 16              |

## ORDER BY

- The order by clause is used to sort the tuples in a query result based on the values of some attributes

### Example

*Question: display the sailors table in the ascending order of sname*

*Query:*

```
SELECT *  
FROM Sailors  
ORDER BY sname
```

*Output:*

| <b>SID</b> | <b>SNAME</b> | <b>RATING</b> | <b>AGE</b> |
|------------|--------------|---------------|------------|
| 32         | Andy         | 8             | 25.5       |
| 85         | Art          | 3             | 25.5       |
| 95         | Bob          | 3             | 63.5       |
| 29         | Brutus       | 1             | 33         |
| 22         | Dustin       | 7             | 45         |
| 64         | Horatio      | 7             | 35         |
| 74         | Horatio      | 9             | 35         |
| 31         | Lubber       | 8             | 55.5       |
| 58         | Rusty        | 10            | 35         |
| 71         | Zorba        | 10            | 16         |



*Question: display the sailors table in the descending order of sname*

*Query:*

```
SELECT *  
FROM Sailors  
ORDER BY sname DESC
```

*Output:*

| <b>SID</b> | <b>SNAME</b> | <b>RATING</b> | <b>AGE</b> |
|------------|--------------|---------------|------------|
| 71         | Zorba        | 10            | 16         |
| 58         | Rusty        | 10            | 35         |
| 31         | Lubber       | 8             | 55.5       |
| 64         | Horatio      | 7             | 35         |
| 74         | Horatio      | 9             | 35         |
| 22         | Dustin       | 7             | 45         |
| 29         | Brutus       | 1             | 33         |
| 95         | Bob          | 3             | 63.5       |
| 85         | Art          | 3             | 25.5       |
| 32         | Andy         | 8             | 25.5       |

## NULL VALUES

- Thus far, we have assumed that column values in a row are always known. In practice column values can be unknown
- We use *null* when the column value is either unknown

### Example

- Insert the row (98,Dan,null,39) to represent Dan into sailors table

*Query:* INSERT INTO Sailors VALUES(98,'Dan',null,39)

*Query:* SELECT \*  
FROM Sailors

*Output:*

| SID | SNAME   | RATING | AGE  |
|-----|---------|--------|------|
| 22  | Dustin  | 7      | 45   |
| 29  | Brutus  | 1      | 33   |
| 31  | Lubber  | 8      | 55.5 |
| 32  | Andy    | 8      | 25.5 |
| 58  | Rusty   | 10     | 35   |
| 64  | Horatio | 7      | 35   |
| 71  | Zorba   | 10     | 16   |
| 74  | Horatio | 9      | 35   |
| 85  | Art     | 3      | 25.5 |
| 95  | Bob     | 3      | 63.5 |
| 98  | Dan     |        | 39   |

## Comparisons Using Null Values

- Consider a comparison such as  $rating = 8$
- *If this is applied to the row for Dan, is this condition true or false? Since Dan's rating is unknown, it is evaluated to the value *unknown**
- This is the case for the comparisons  $rating > 8$  and  $rating < 8$  as well
- SQL also provides a special comparison operator IS NULL to test whether a column value is *null*

- for example, we can say rating IS NULL, which would evaluate to true on the row representing Dan
- We can also say rating IS NOT NULL, which would evaluate to false on the row for Dan

### Example

#### *Query:*

```
SELECT *  
FROM sailors  
WHERE rating IS NULL
```

#### *Output:*

| SID | SNAME | RATING | AGE |
|-----|-------|--------|-----|
| 98  | Dan   |        | 39  |

## Logical Connectives AND, OR, and NOT

- what about Boolean expressions such as

*rating = 8 OR age < 40*

*rating = 8 AND age < 40?*

- *Considering the row for Dan again, because age < 40, the first expression evaluates to true regardless of the value of rating, but what about the second? We can only say unknown*
- The expression NOT unknown is defined to be unknown

▪ OR of two arguments evaluates to *true* if either argument evaluates to *true*, and to *unknown* if one argument evaluates to *false* and the other evaluates to *unknown*

| A     | B     | A OR B |
|-------|-------|--------|
| T     | F (U) | T      |
| F (U) | T     | T      |
| F     | U     | U      |
| U     | F     | U      |

▪ AND of two arguments evaluates to *false* if either argument evaluates to *false*, and to *unknown* if one argument evaluates to *unknown* and the other evaluates to *true* or *unknown*

| A     | B     | A AND B |
|-------|-------|---------|
| T     | T     | T       |
| F     | U (T) | F       |
| U (T) | F     | F       |
| U     | U (T) | U       |
| U (T) | U     | U       |



## Impact on SQL Constructs

- In the presence of *null values*, any row that evaluates to false or to unknown is eliminated
- If we compare two *null values* using  $=$ , the result is unknown! In the context of duplicates, this comparison is implicitly treated as true, which is an anomaly
- The arithmetic operations  $+$ ,  $-$ ,  $*$ ,  $/$  and  $=$  all return *null* if one of their arguments is null

### Example

*Query:*

```
SELECT sid, rating, sid+rating  
FROM Sailors
```

*Output:*

| SID | RATING | SID + RATING |
|-----|--------|--------------|
| 22  | 7      | 29           |
| 29  | 1      | 30           |
| 31  | 8      | 39           |
| 32  | 8      | 40           |
| 58  | 10     | 68           |
| 64  | 7      | 71           |
| 71  | 10     | 81           |
| 74  | 9      | 83           |
| 85  | 3      | 88           |
| 95  | 3      | 98           |
| 98  | -      | -            |

- nulls can cause some unexpected behavior with aggregate operators
- COUNT(\*) handles *null values just like other values*, that is, they get counted

### Example

*Query:*


```
SELECT COUNT(*)  
FROM Sailors
```

*Output:*

| COUNT(*) |
|----------|
| 11       |

- All the other aggregate operators (COUNT, SUM, AVG, MIN, MAX, and variations using DISTINCT) simply discard *null values*

## Outer Joins

- join operation that rely on *null values*, called **outer joins**
- Consider the join of two tables, say Sailors  Reserves
- In a *full outer join*, ‘matching rows’ plus ‘Sailors rows without a matching Reserves rows’ (columns inherited from Reserves assigned *null* values) plus ‘Reserves rows without a matching Sailors rows’ (columns inherited from Sailors assigned *null* values) appear in the result

- In a *left outer join*, ‘matching rows’ plus ‘Sailors rows without a matching Reserves rows’ (columns inherited from Reserves assigned *null* values) appear in the result
- In a *right outer join*, ‘matching rows’ plus ‘Reserves rows without a matching Sailors rows’ (columns inherited from Sailors assigned *null* values) appear in the result
- *Note:* In *inner join* only matching rows appear in the result

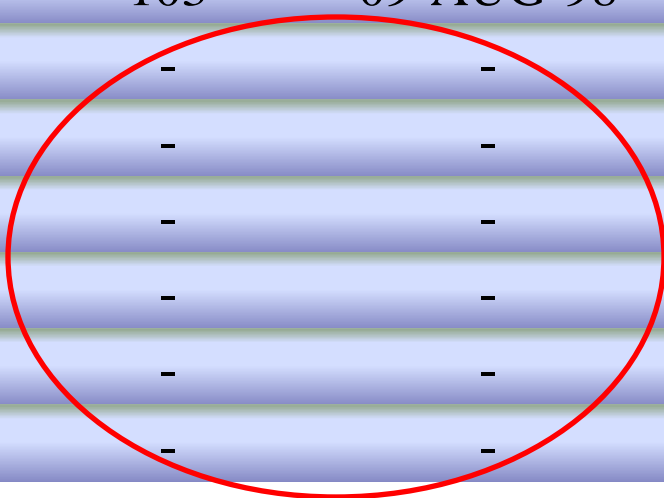
## Example

### Query:

```
SELECT S.sid,S.sname,R.bid,R.day  
FROM Sailors S LEFT OUTER JOIN Reserves R ON  
S.sid=R.sid
```

*Output:*

| <b>SID</b> | <b>SNAME</b> | <b>BID</b> | <b>DAY</b> |
|------------|--------------|------------|------------|
| 22         | Dustin       | 101        | 10-OCT-98  |
| 22         | Dustin       | 102        | 10-OCT-98  |
| 22         | Dustin       | 103        | 10-AUG-98  |
| 22         | Dustin       | 104        | 10-JUL-98  |
| 31         | Lubber       | 102        | 11-OCT-98  |
| 31         | Lubber       | 103        | 11-JUN-98  |
| 31         | Lubber       | 104        | 11-DEC-98  |
| 64         | Horatio      | 101        | 09-MAY-98  |
| 64         | Horatio      | 102        | 09-AUG-98  |
| 74         | Horatio      | 103        | 09-AUG-98  |
| 71         | Zorba        | -          | -          |
| 85         | Art          | -          | -          |
| 58         | Rusty        | -          | -          |
| 32         | Andy         | -          | -          |
| 29         | Brutus       | -          | -          |
| 95         | Bob          | -          | -          |



## Disallowing Null Values

- We can disallow *null* values by specifying NOT NULL as part of the field definition, for example,

*sname* VARCHAR2(20) NOT NULL

- The fields in a primary key are not allowed to take on *null* values
- There is an implicit NOT NULL constraint for every field listed in a PRIMARY KEY constraint